

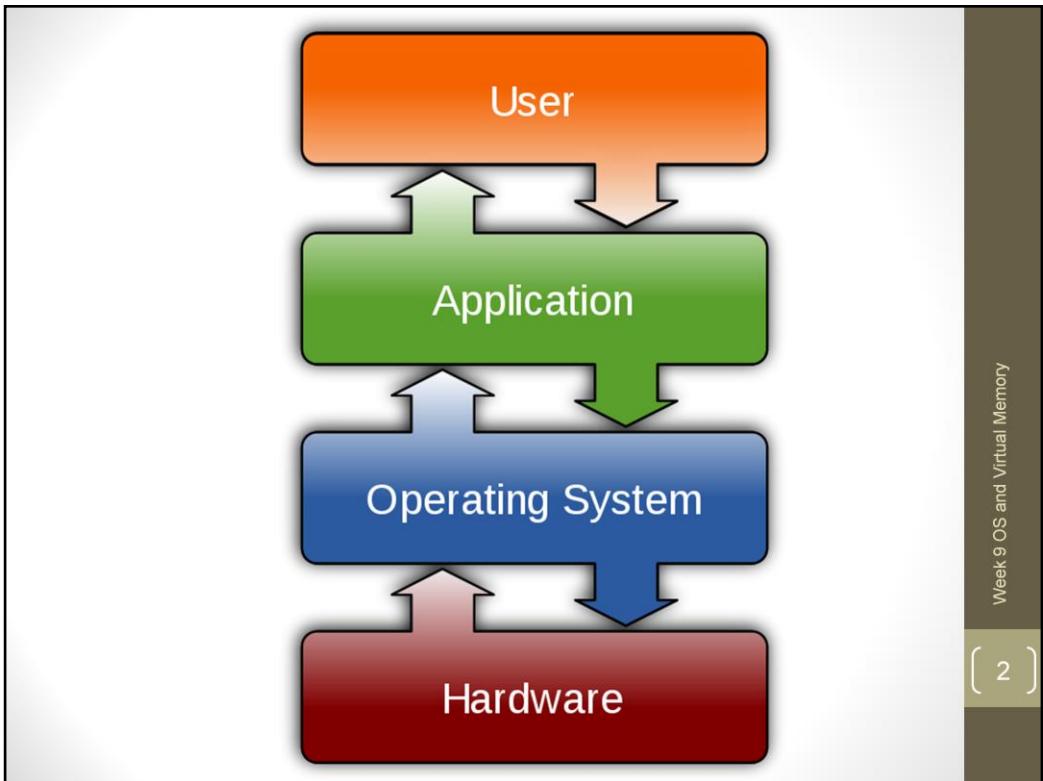
Computing Systems

Week 9 OS and Virtual Memory

Operating Systems & Virtual Memory

[1]





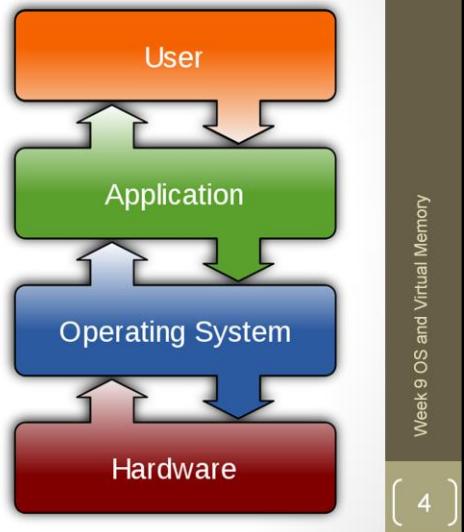
Operating Systems (OS)

- OS is a collection of software that manages computer hardware resources and provides common services for computer programs.
- The operating system is an essential component of the system software in a computer system.
- Application programs usually require an operating system to function.

(3)

Operating Systems

Software that sits between applications (e.g. office applications or games) and the hardware itself



Features and Components

- Kernel Space/User Space
- Process management
- Interrupts & I/O
- Memory management
- File system
- Device drivers
- Networking (TCP/IP, UDP)
- Security (Process/Memory protection)

[5]

List taken from

http://en.wikipedia.org/wiki/Operating_system

CC-BY-SA

Popular Operating Systems

- Windows
- *nix (**Unix** family of operating systems
 - including **Mac OS X, Linux, iOS, Android** and many others)
- **DOS** (pre-Windows Operating Systems for PCs)
- Also: **Symbian, WebOS, BlackBerry OS, ...**

```
Current date is Tue 1-01-1980
Enter new date:
Current time is 7:48:27.13
Enter new time:

The IBM Personal Computer DOS
Version 1.10 (C)Copyright IBM Corp 1981, 1982

A>dir/w
COMMAND COM   FORMAT COM   CHDKSK COM   SYS     COM   DISKCOPY COM
DISKCOMP COM   COMP    COM   EXE2BIN EXE   MODE    COM   EDLIN   COM
DEBUG    COM   LINK    EXE   BASIC   COM   BASICA   COM   ART     BAS
SAMPLES BAS   MORTGAGE BAS   COLORBAR BAS   CALENDAR BAS   MUSIC    BAS
DONKEY  BAS   CIRCLE  BAS   PIECHART BAS   SPACE    BAS   BALL    BAS
COMM    BAS

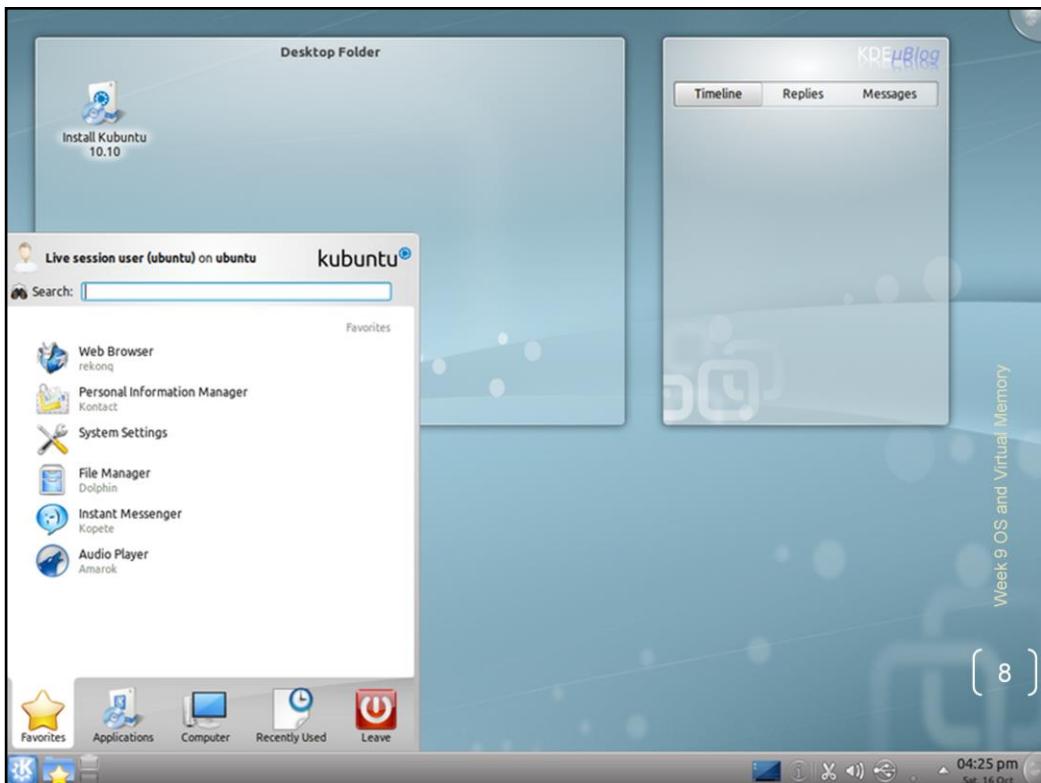
      26 File(s)
A>dir command.com
COMMAND COM   4959   5-07-82  12:00p
      1 File(s)
A>
```

Week 9 OS and Virtual Memory

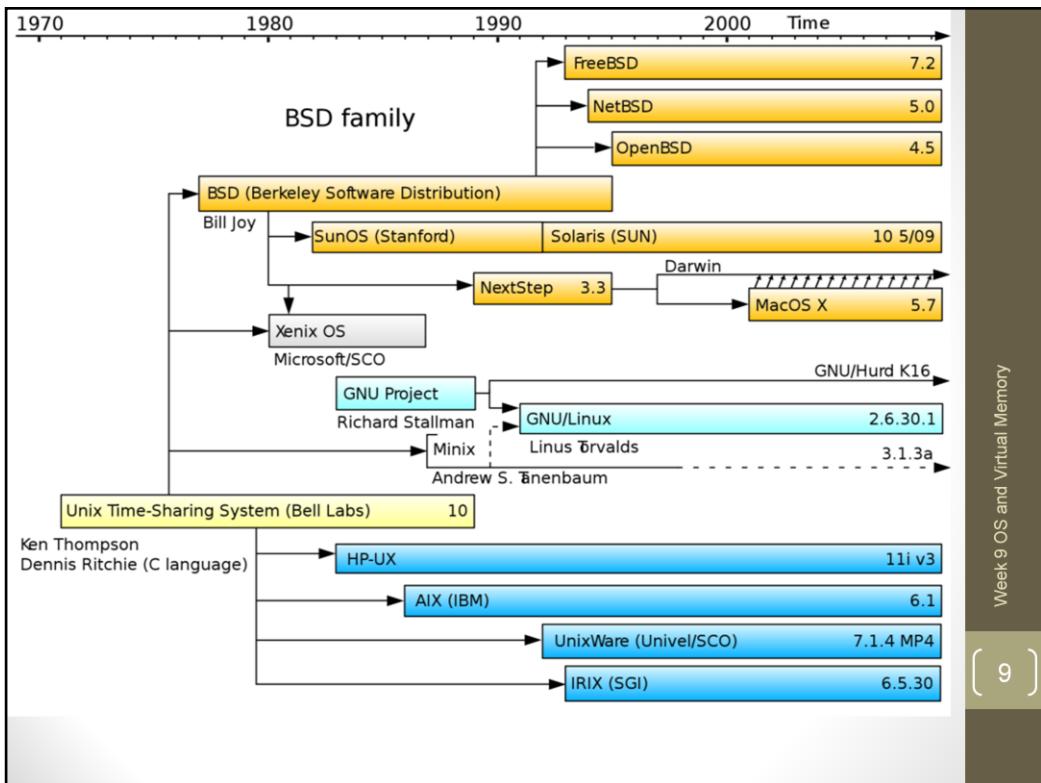
(7)

This is PC-Dos – the IBM version of MS-DOS, the same DOS as sold by Microsoft, but with the IBM name instead.

Other companies sold compatible DOS systems for PCs, but Microsoft's licensing agreement with IBM allowed Microsoft to become the dominant software company for Operating Systems – and to exploit this to grow into other areas of computing.



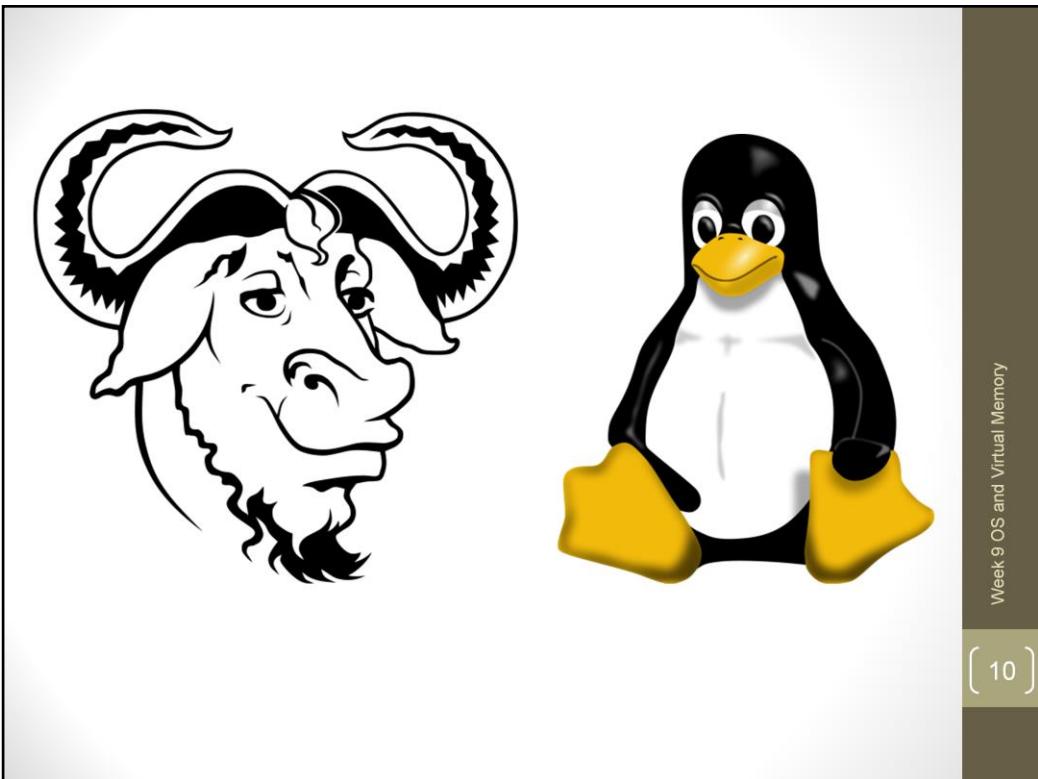
This image shows an Ubuntu KDE desktop (Kubuntu). Ubuntu is a popular Linux distribution. Linux is open source, and the license allows companies or individuals to create their own distributions or versions. The Linux OS can work with (or without!) a GUI component. The X Windows System (aka X11) provides the GUI API, but two distinct front-ends desktop environments with their own sets of applications have emerged: KDE and Gnome. (Prior to KDE and Gnome, *nix applications had generally inconsistent GUIs and often looked visually poor in comparison to e.g. Windows or Mac OS)



The UNIX family has a long and convoluted history.

A multi-user OS underdevelopment by Bell Labs, MIT and GE ‘Multics’ was innovative but problematic. Bell Labs tried again. A team including Ken Thompson, Dennis Ritchie and Brian Kernighan developed *Unics* – then renamed Unix. Along the way, Ritchie and Kernighan created the C programming language to write the new OS.

Overtime many versions of UNIX were developed, with many major IT companies having their own version – including Microsoft.

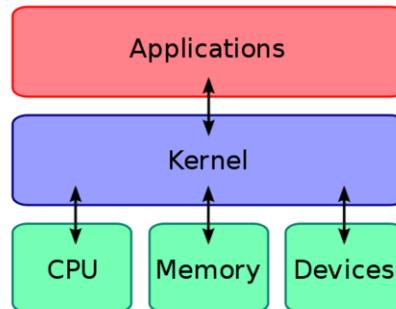


GNU (GNU is Not Unix), Tux the Penguin (Linux mascot)

The GNU project started in 1983 by Richard Stallman with the goal of creating a complete free and Open Source version of the UNIX operating system. Over the following years the GNU project was able to make substantial progress with the new OS, creating much of the software required – compilers, editors, libraries and utilities. But progress was slow in developing the OS Kernel – the core of the new OS. In 1991 Linus Torvalds wrote and released Linux – an OS core that could be used with GNU to create a complete and usable OS.

The Kernel

- Most basic set of OS functionality exists in the **OS Kernel**
- **Kernel** has some key tasks:
 - **Manage different running processes**
 - **Allocate memory**
 - **Process I/O**
- Kernel may also perform **file management**, run **device drivers**, and carry out a range of other tasks



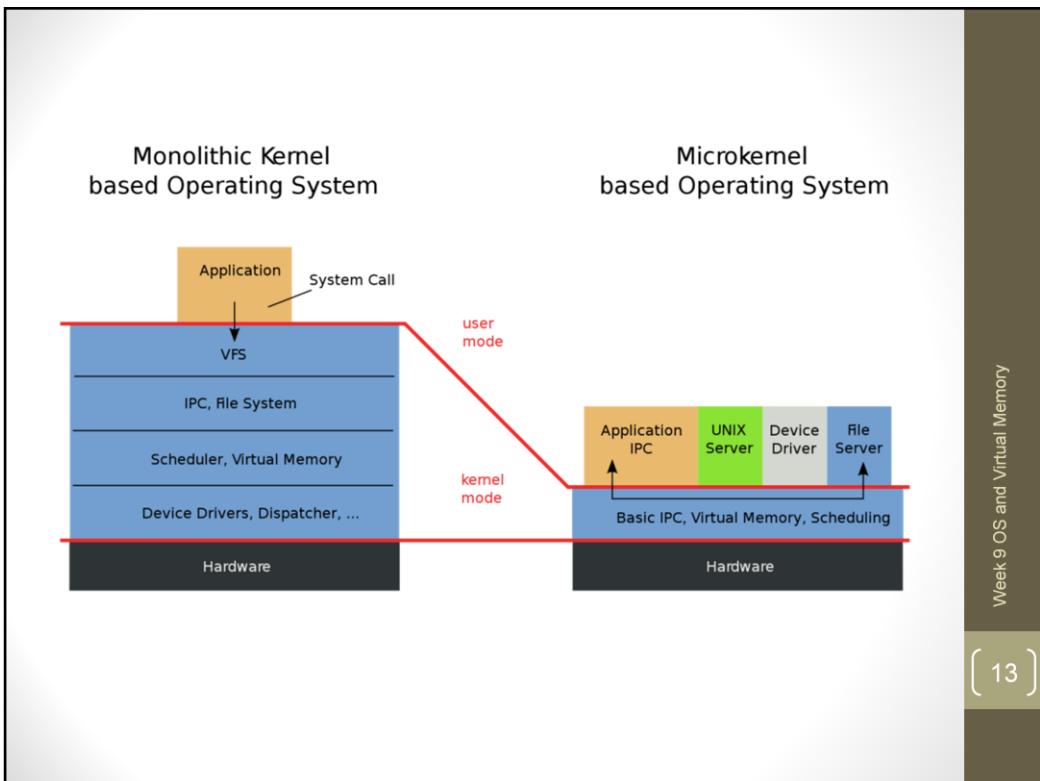
Week 9 OS and Virtual Memory

[11]

Linux actually refers to the Kernel – GUI and many other elements are not part of the core of Linux. A Linux distribution is made up (typically) of the Linux Kernel, a windowing system, desktop environment and a range of applications. The kernel itself is much more limited... Which is why some Linux systems look very different to one another.

Kernel

- The **kernel** is a computer program that ***manages*** **input/output** requests from **software** and ***translates*** them into **data processing instructions** for the **central processing unit** and other electronic **components** of a computer.
- It ***connects*** the **application software** to the **hardware** of a computer.
- It ***manages memory*** access for programs in the **RAM**, it determines which **programs** get ***access*** to which **hardware** resources, it ***sets up*** or ***resets*** the **CPU's operating states** for optimal operation at all times, and it ***organizes*** the **data** for **long-term non-volatile** storage with **file systems** on such **media** as **disks, tapes, flash** memory, etc.



Windows and **Linux** both use a fairly **monolithic** kernel design.

Micro-kernels move parts of the **operating system** into services that run **outside** of the kernel. This can mean that some failures (in e.g. The file server or with device drivers) are less likely to cause crashes, and the system may be more secure. The cost is that there may be a performance hit.

IPC = “**Inter Process Communication**”

VFS = “**Virtual File System**” – more later...

User Space (User Mode)

- **Kernel Space** has direct access to the *memory* used by the *operating system* itself
 - **Security** issues exist in making this accessible to all applications
- **Applications** and **OS** modules may exist in a **User Space**
 - Limits direct access to the *memory* of other *applications* and to *hardware*
 - Forces programs to use the *kernel API* interface

User space provides some protection to the system from malicious or poorly written programs that might crash or damage the system.

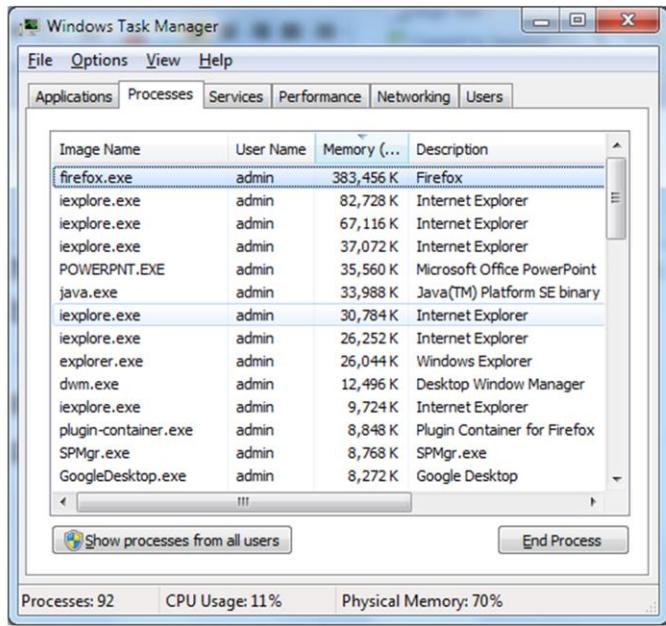
A little more on this later!

Process Management

- Modern **OS** allow *multiple programs* to be run by a *computer* at *one time*
 - Even if there is only *one processor*
- Requires **OS** to switch between *running programs/processes*
 - If one process is *waiting* for user input, *pause* it and run a *different* process
- Some process may need to interact
 - **OS** manages ***Inter-Process Communication (IPC)***

Multi-Processing & Multi-Tasking

- ***Multi-Processing*** uses **multiple processors** (or **processor cores**) to run **multiple tasks** (programs)
- ***Multi-Tasking*** allows **multiple tasks** to run on **one or more processors**
 - May have **multiple tasks** per **processor**
 - **Rapid switching** between **tasks**, performing a little bit of each, allows computer to give the appearance of running all the tasks at the same time



Three State Process Management

- Process ***waiting*** for I/O (or a signal from another process) is **BLOCKED**
- **States and transitions:**
 - **RUNNING:** The process that is currently being *executed*
 - **READY:** A process that is *queuing* and prepared to execute when given the opportunity
 - **BLOCKED:** A process that *cannot execute* until some event occurs, such as the *completion* of an I/O operation

[18]

States and transitions text from
[\(CC-BY-SA\)](http://en.wikipedia.org/wiki/Process_management_(computing))

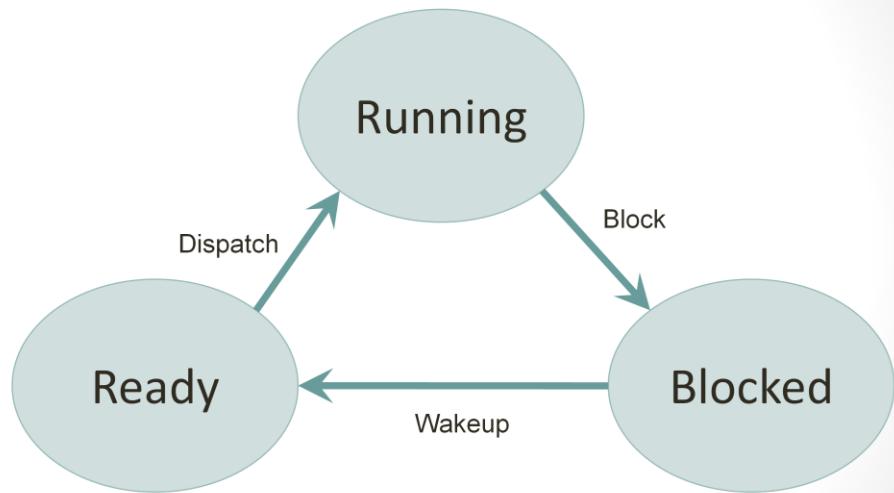
Transitions (again, text from Wikipedia, as above):

Processes entering the system must go initially into the **READY** state, processes can only enter the **RUNNING** state via the **READY** state. Processes normally leave the system from the **RUNNING** state. For each of the three states, the process occupies space in main memory. While the reason for most transitions from one state to another might be obvious, some may not be so clear.

RUNNING → READY The most common reason for this transition is that the running process has reached the maximum allowable time for uninterrupted execution; i.e. time-out occurs. Other reasons can be the imposition of priority levels as determined by the [scheduling](#) policy used for the Low Level [Scheduler](#), and the arrival of a higher priority process into the **READY** state.

RUNNING → BLOCKED A process is put into the **BLOCKED** state if it requests something for which it must wait. A request to the OS is usually in the form of a system call, (i.e. a call from the running process to a function that is part of the OS code). For example, requesting a file from disk or a saving a section of code or data from memory to a file on disk.

Process Transitions



Week 9 OS and Virtual Memory

[19]

With **pre-emptive** multitasking, the OS will regularly pause or block running processes to allow other processes to run. **Co-operative** multitasking requires process to voluntarily suspend or block themselves. These is also a five state variant of this diagram. From wikipedia

[http://en.wikipedia.org/wiki/Process_management_\(computing\) \(CC-BY-SA\) :](http://en.wikipedia.org/wiki/Process_management_(computing) (CC-BY-SA) :)

Five-state process management model

While the three state model is sufficient to describe the behaviour of processes with the given events, we have to extend the model to allow for other possible events, and for more sophisticated design. In particular, the use of a portion of the hard disk to emulate main memory (so called virtual memory) requires additional states to describe the state of processes which are suspended from main memory, and placed in virtual memory (on disk). Of course, such processes can, at a future time, be resumed by being transferred back into main memory. The Medium Level Scheduler controls these events. A process can be suspended from the *RUNNING*, *READY* or *BLOCKED* state, giving rise to two other states, namely, *READY SUSPEND* and *BLOCKED SUSPEND*. A *RUNNING*

process that is suspended becomes *READY SUSPEND*, and a *BLOCKED* process that is suspended becomes *BLOCKED SUSPEND*. A process can be suspended for a number of reasons; the most significant of which arises from the process being swapped out of memory by the memory management system in order to free memory for other processes. Other common reasons for a process being suspended are when one suspends execution while debugging a program, or when the system is monitoring processes. For the five-state process management model, consider the following transitions described in the next sections.

BLOCKED → BLOCKED SUSPEND If a process in the *RUNNING* state requires more memory, then at least one *BLOCKED* process can be swapped out of memory onto disk. The transition can also be made for the *BLOCKED* process if there are *READY* processes available, and the OS determines that the *READY* process that it would like to dispatch requires more main memory to maintain adequate performance.

BLOCKED SUSPEND → READY SUSPEND A process in the *BLOCKED SUSPEND* state is moved to the *READY SUSPEND* state when the event for

which it has been waiting occurs. Note that this requires that the state information concerning suspended processes be accessible to the OS.

READY SUSPEND → READY When there are no *READY* processes in main memory, the OS will need to bring one in to continue execution. In addition, it might be the case that a process in the *READY SUSPEND* state has higher priority than any of the processes in the *READY* state. In that case, the OS designer may dictate that it is more important to get at the higher priority process than to minimise swapping.

READY → READY SUSPEND Normally, the OS would be designed so that the preference would be to suspend a *BLOCKED* process rather than a *READY* one. This is because the *READY* process can be executed as soon as the [CPU](#) becomes available for it, whereas the *BLOCKED* process is taking up main memory space and cannot be executed since it is waiting on some other event to occur. However, it may be necessary to suspend a *READY* process if that is the only way to free a sufficiently large block of main memory. Finally, the OS may choose to suspend a lower-priority *READY* process rather than a higher-priority *BLOCKED* process if it believes that the

BLOCKED process will be ready soon.

Multi-Threading: Processes and Threads

- Each **process** has its own ***private memory space***
 - When switching between **processes**, the ***memory context*** also needs ***switched***
- **Threads** work similarly to **processes**, except **threads** can ***share a memory context***
 - Switching between **threads** in a **process** is ***faster***
 - ***Shared access*** to ***data*** across **threads** rather than needing to ***pass*** data between **processes**

[20]

Multi-Threading: The ability of an operating system to **execute** different **parts** of a **program**, called **threads, simultaneously**. The programmer must carefully design the program in such a way that all the threads can run at the same time without interfering with each other.

Multiprocessing: Refers to a computer system's ability to support more than one **process (program)** at the same time. Multiprocessing operating systems enable **several programs** to ***run concurrently***. **UNIX** is one of the most widely used multiprocessing systems, but there are many others, including OS/2 for high-end PCs. Multiprocessing systems are much more complicated than single-process systems because the operating system must allocate resources to competing processes in a reasonable manner.

I/O and Interrupts

- Already looked at **interrupts** and **I/O**
- **Operating System** responsible for taking **inputs** and passing these onto user applications where necessary
 - E.g. Passing mouse click or key press to the current application

[21]

Input/Output (I/O)

- I/O is the communication between an information processing system (such as a computer) and the outside world, possibly a human or another information processing system.
- **Inputs** are the **signals** or **data received** by the **system**, and **outputs** are the **signals** or **data sent** from it.

[22]

Interrupt

- An **interrupt** is a signal to the processor emitted by hardware or software indicating an event that needs immediate attention.
- An **interrupt** alerts the processor to a **high-priority** condition requiring the *interruption* of the current **code** the processor is **executing** (the current *thread*).

[23]

Interrupt

- The processor responds by suspending its current activities, saving its state, and executing a small program called an ***interrupt handler*** (or ***interrupt service routine, ISR***) to deal with the event.
- This *interruption* is *temporary*, and after the interrupt handler *finishes*, the processor *resumes execution* of the *previous thread*.

[24]

Device Drivers

- **Programs** written to work with specific hardware systems – and to provide a simpler, abstract, **interface to the OS and applications**
 - Kernel space and user space drivers
- **OS** and application developers should **not need** to **know** the specific **low-level details** of all the different **hardware** systems that might be used
 - Cannot know about **future devices**

[25]

Memory Management

- **Memory management** is the act of managing computer memory.
- The essential requirement of memory management is to provide ways to **dynamically allocate portions of memory to programs** at their request, and free it for reuse when no longer needed.

[26]

Main Memory

- **Memory address space** gives a **unique numerical address** to every **byte**
 - Or **set of bytes**
- **Addresses start at zero**
- Possible size of address space depends on computer design
 - E.g. **32 bit x86** CPUs limited to **4 GB** physical **memory**

5FFF	...
5FFE	A
5FFD	T
5FFC	A
	D
	...
0004	E
0003	M
0002	O
0001	S
0000	

Week 9 OS and Virtual Memory
[27]

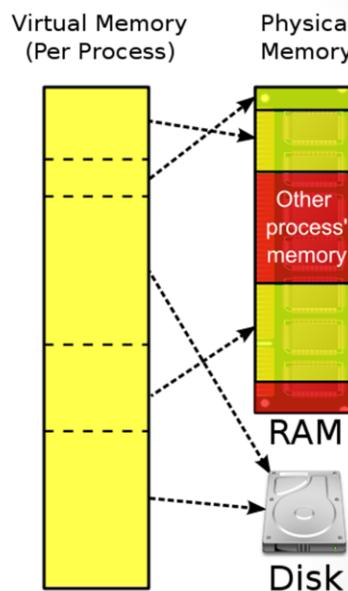
The first computers and first home computers had minimal OS memory management. Only one program could be run at a time, and it would have access to all of the computers main memory. In the 1970's and 80's this was often used (and abused) by expert programmers to squeeze the maximum performance possible from the computers of the time. Parts of main memory were also given over to system I/O – to the extent that by writing data values to specific addresses of main memory, output could be drawn to the screen, or locations in memory directly read to obtain input.

Memory Pages & Segments

- Instead of using physical address of each byte, can use a system where a memory address is stored as a reference to a **page** or **segment** of memory
 - Plus **offset** to data in the block
- Early architectures had very tight limits on the address space available to individual programs (**64K** on early **PCs**)
 - **Segment:** Offset addressing increased total available memory

Virtual Memory

- With limited **physical memory**, *multi-process OS* could struggle to fit all programs into memory
- **Virtual Memory** allows the use of secondary **storage** to give the appearance of more memory than actually exists



Week 9 OS and Virtual Memory

[29]

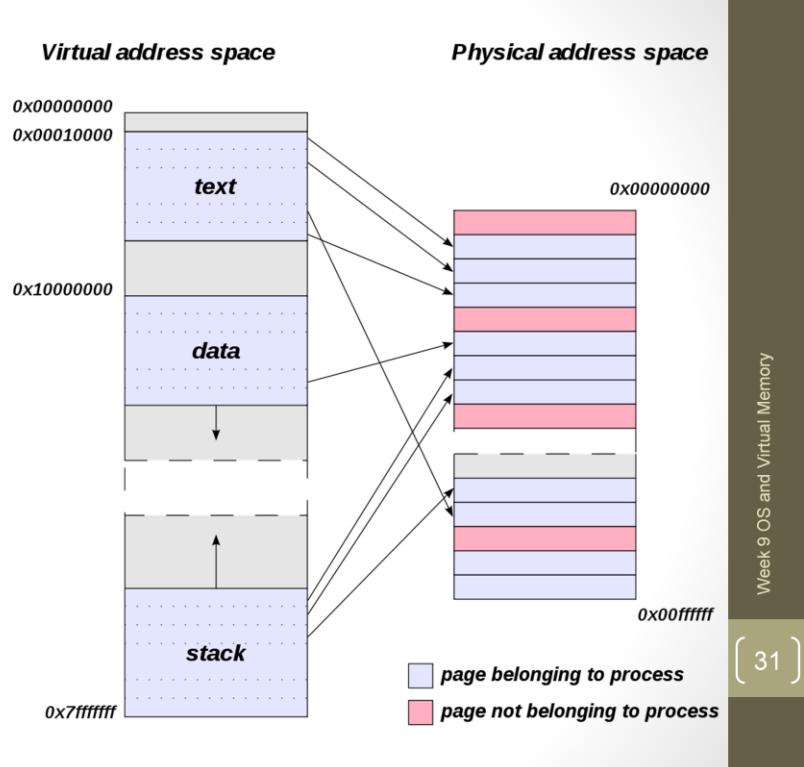
Virtual Memory

- On disk *swap files* or *swap partitions* used to add *virtual memory* to *physical memory*
- Programs may have *partial residence*
 - Partly in physical **memory**, partly in VM
- *Scatter loading* allows different programs to exist on pages scattered through the physical and virtual memory
- **Scatter loading** and **partial residence** is invisible to the applications themselves

[30]

Page Table

Table maps virtual addresses to physical addresses



From: http://en.wikipedia.org/wiki/Page_table

Relationship between pages addressed by virtual addresses and the frames in physical memory, within a simple address space scheme. Physical memory can contain pages belonging to many processes. Pages can be swapped to disk if used infrequently, or if physical memory is full. Not all pages are in physical memory in the above diagram.

Virtual Memory Operation

- Application attempts to access a **VM address**
- **OS** checks virtual memory address on page table
 - If page currently in memory (a page '**hit**'), the **physical address is returned**
 - Else (a page '**miss**') the page needs **loaded** into **physical memory**, perhaps replacing an existing page which gets copied out to disk (**swap**)
- **Thrashing** can occur when low physical memory requires constant disk swapping

[32]

Virtual Memory & User Space

- Applications running in *user space* have their own *virtual memory address spaces*
- When they attempt to *read/write* memory, virtual addresses are *translated* to *physical ones* by the **OS** & the **hardware memory management unit**
- Attempts to access discs, etc., similarly mediated by OS
- This is all largely *invisible* to the *applications*

[33]

File System

- A **file system** is used to **control** how **data** is *stored* and *retrieved*.
- Without a file system, information placed in a storage area would be one large body of data with no way to tell where one piece of information stops and the next begins.
- By separating the data into *individual* pieces, and giving each piece a *name*, the information is easily *separated* and *identified*.
- Taking its name from the way paper-based information systems are named, each piece of *data* is called a "*file*".
- The *structure* and *logic* rules used to manage the *groups* of *information* and their *names* is called a "**file system**".

File System

- Modern file systems allow creation of folders, subfolders, and long file names.
 - *nix file systems allow spaces in file names
 - – Windows did not until Windows 95.

Networking

- OS will typically include a subsystem to support **networking**.
- This is part of **kernel** in **Windows**.
- In computer **networks**, **networked computing devices** pass ***data*** to each other along ***data connections***.
- The connections (***network links***) between **nodes** are established using either **cable media** or **wireless media**.

[36]

Networking

- **Nodes** can include **hosts** such as **personal computers**, phones, **servers** as well as **networking hardware**.
- Two such **devices** are said to be **networked together** when one **device** is able to **exchange information** with the **other device**, whether or **not** they have a **direct connection** to each other.

Security & Reliability

- Applied to both computers and computer networks.
 - The field covers all the **processes** and **mechanisms** by which computer-based **equipment, information** and **services** are **protected** from **unintended** or **unauthorized** access, change or destruction.
- Additional OS support to limit the potential damage caused by **malicious** or **poorly written applications** and **drivers**.

Quick quiz

- Join the ‘Socrative’ app ‘Room 642124’ and try the quick quiz.

Further Reading

- Wikipedia:
 - Operating Systems, memory management, paging, virtual memory, process management, Unix, Linux, etc.
- PCH
 - Chapter 13
- Modern Operating Systems by Andrew Tanenbaum

[40]

Required Reading For Next Week

- HCW:
 - How the Processor Uses Registers (p66-67)
 - How Multi-core processors work (p72-73)
- Online Microprocessor tutorial at:
 - <http://www.eastaughs.fsnet.co.uk/cpu/index.htm>
 - Read: Introduction, CPU Structure (all), Instruction Execution (Introduction, Instruction Sets, Execution Cycle)
- Additional material/links may be posted on Blackboard: Please check!



This work by Daniel Livingstone at the University of the West of Scotland is licensed under the Creative Commons Attribution-ShareAlike 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/3.0/> or send a letter to Creative Commons, 171 Second Street, Suite 300, San Francisco, California, 94105, USA.

Image Credits:

Title Image – CC-BY-SA Daniel Livingstone, 2011

Operating System placement diagram CC-BY-SA Golftheimer, http://en.wikipedia.org/wiki/File:Operating_system_placement.svg

Kubuntu screenshot is free software; you can redistribute it and/or modify it under the terms of the **GNU General Public License** as published by the Free Software Foundation; either version 2 of the License, or any later version. This work is distributed in the hope that it will be useful, but **without any warranty**; without even the implied warranty of **merchantability** or **fitness for a particular purpose**. See [version 2](#) and [version 3 of the GNU General Public License](#) for more details

Simplified UNIX family tree, Public Domain, from http://en.wikipedia.org/wiki/File:Unix_history.svg

GNU Gnu by Aurelio Heckert. Copyleft: This work of art is free; you can redistribute it and/or modify it according to terms of the [Free Art License](#). You will find a specimen of this license on the [Copyleft Attitude site](#) as well as on other sites

Tux Penguin: The original [Tux](#), the official [Linux](#) mascot created by Larry Ewing [1] in 1996: *Permission to use and/or modify this image is granted provided you acknowledge me lewing@isc.tamu.edu and The GIMP if someone asks.*

Kernel schematic, CC-BY-SA Bobbo, http://en.wikipedia.org/wiki/File:Kernel_layout.svg

Monolithic vs Micro Kernel schematic, Public Domain by Woottoo <http://en.wikipedia.org/wiki/File:OS-structure.svg>

Windows task manager is (c) Microsoft. User here under fair use provisions to illustrate range of processes on a modern Windows PC. Screenshot taken by Daniel Livingstone.

Virtual Memory schematic CC-BY-SA Ehamberg, http://en.wikipedia.org/wiki/File:Virtual_memory.svg

Page Table diagram, Copyright © 2011 [en:User:Dysprosia](#): Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met: 1. Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer. 2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution. 3. Neither the name of [en:User:Dysprosia](#) nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.
http://en.wikipedia.org/wiki/File:Virtual_address_space_and_physical_address_space_relationship.svg

Product names, logos, brands, and other trademarks featured or referred to within the these notes are the property of their respective trademark holders. These notes have been produced without affiliation, sponsorship or endorsement from the trademark holders.